

Student Name : Sagar Kapase
Roll No : BEA-7

Group B Machine Learning

Assignment 4 K Means Clustering

K Means Clustering

```
In [11]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
```


```
In [4]: df=pd.read_csv('sales_data_sample.csv',encoding='latin-1')
```

```
In [5]: df.sample(5)
```

```
Out[5]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE
1045	10301	23	100.00	9	4011.66	10/5/2005
2498	10308	21	100.00	12	2224.95	10/15/2005
2275	10413	24	49.71	6	1193.04	5/5/2005
2485	10133	24	77.64	8	1863.36	6/27/2005
428	10194	21	93.34	10	1960.14	11/25/2005

5 rows × 25 columns



```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ORDERNUMBER           2823 non-null  int64  
1   QUANTITYORDERED       2823 non-null  int64  
2   PRICEEACH             2823 non-null  float64 
3   ORDERLINENUMBER       2823 non-null  int64  
4   SALES                 2823 non-null  float64 
5   ORDERDATE             2823 non-null  object  
6   STATUS                2823 non-null  object  
7   QTR_ID                2823 non-null  int64  
8   MONTH_ID              2823 non-null  int64  
9   YEAR_ID               2823 non-null  int64  
10  PRODUCTLINE           2823 non-null  object  
11  MSRP                  2823 non-null  int64  
12  PRODUCTCODE           2823 non-null  object  
13  CUSTOMERNAME          2823 non-null  object  
14  PHONE                 2823 non-null  object  
15  ADDRESSLINE1          2823 non-null  object  
16  ADDRESSLINE2          302 non-null   object  
17  CITY                  2823 non-null  object  
18  STATE                 1337 non-null  object  
19  POSTALCODE            2747 non-null  object  
20  COUNTRY               2823 non-null  object  
21  TERRITORY             1749 non-null  object  
22  CONTACTLASTNAME       2823 non-null  object  
23  CONTACTFIRSTNAME      2823 non-null  object  
24  DEALSIZE              2823 non-null  object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: ORDERNUMBER      0
        QUANTITYORDERED  0
        PRICEEACH        0
        ORDERLINENUMBER  0
        SALES             0
        ORDERDATE        0
        STATUS           0
        QTR_ID           0
        MONTH_ID         0
        YEAR_ID          0
        PRODUCTLINE      0
        MSRP             0
        PRODUCTCODE      0
        CUSTOMERNAME     0
        PHONE            0
        ADDRESSLINE1     0
        ADDRESSLINE2     2521
        CITY             0
        STATE            1486
        POSTALCODE       76
        COUNTRY          0
        TERRITORY        1074
        CONTACTLASTNAME  0
        CONTACTFIRSTNAME 0
        DEALSIZE         0
        dtype: int64
```

```
In [8]: df.shape
```

```
Out[8]: (2823, 25)
```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 0
```

In [10]: `df.corr()`

Out[10]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP
ORDERNUMBER	1.000000	0.065543	-0.002935	-0.055550	0.039919	-0.051383	-0.039723	0.904596	-0.010280
QUANTITYORDERED	0.065543	1.000000	0.005564	-0.018397	0.551426	-0.035323	-0.039048	0.069535	0.017881
PRICEEACH	-0.002935	0.005564	1.000000	-0.020965	0.657841	0.008712	0.005152	-0.005938	0.670625
ORDERLINENUMBER	-0.055550	-0.018397	-0.020965	1.000000	-0.058400	0.040716	0.034016	-0.057367	-0.021067
SALES	0.039919	0.551426	0.657841	-0.058400	1.000000	0.000000	0.000000	0.000000	0.000000
QTR_ID	-0.051383	-0.035323	0.008712	0.040716	0.000000	1.000000	0.000000	0.000000	0.000000
MONTH_ID	-0.039723	-0.039048	0.005152	0.034016	0.000000	0.000000	1.000000	0.000000	0.000000
YEAR_ID	0.904596	0.069535	-0.005938	-0.057367	0.000000	0.000000	0.000000	1.000000	0.000000
MSRP	-0.010280	0.017881	0.670625	-0.021067	0.000000	0.000000	0.000000	0.000000	1.000000

In [14]: `df['CITY'].value_counts()`

Out[14]:

Madrid	304
San Rafael	180
NYC	152
Singapore	79
Paris	70
...	
Graz	15
Los Angeles	14
Munich	14
Burbank	13
Charleroi	8

Name: CITY, Length: 73, dtype: int64

In [16]: `df['STATE'].value_counts().count()`

Out[16]: 16

In [18]: `df1=df.select_dtypes(exclude='object')`

In [19]: `df1.shape`

Out[19]: (2823, 9)

In [20]: `df1.sample(5)`

Out[20]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID
1307	10315	36	100.00	7	3602.16	4
2222	10104	35	47.62	11	1666.70	1
2500	10328	27	100.00	8	2762.10	4
1762	10328	48	58.92	1	2828.16	4
2242	10335	40	60.60	3	2424.00	4

In [21]: `std_scaler= StandardScaler()
df_scaled = std_scaler.fit_transform(df1)`

In [23]: `df_scaled`

Out[23]:

```
array([[ -1.64794709, -0.52289086,  0.5969775 , ..., -1.39290889,
        -1.16517009, -0.14224584],
       [ -1.4958875 , -0.11220131, -0.11445035, ..., -0.57233673,
        -1.16517009, -0.14224584],
       [ -1.35468931,  0.60650538,  0.54938372, ..., -0.02528862,
        -1.16517009, -0.14224584],
       ...,
       [  1.38238338,  0.81185016,  0.81015797, ..., -1.11938483,
         1.69382614, -1.16263387],
       [  1.50185877, -0.11220131, -1.06186404, ..., -1.11938483,
         1.69382614, -1.16263387],
       [  1.68650256,  1.2225397 , -0.89925195, ..., -0.57233673,
         1.69382614, -1.16263387]])
```

In [24]: `df2=pd.DataFrame(df_scaled,columns=df1.columns)`

In [26]: `df2.sample(5)`

Out[26]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_I
1136	1.197740	1.017195	-0.524451	-0.583696	-0.144058	-1.42703
386	0.708977	-0.933580	-1.253231	0.126347	-1.105602	1.06535
1201	-0.073044	-0.420218	-1.527886	-0.347015	-1.040352	-0.59624
86	-0.855065	0.606505	0.810158	1.073072	2.899212	1.06535
1101	-0.225104	0.606505	-0.519989	-0.110334	-0.300791	-0.59624

In [25]: `df2.corr()`

Out[25]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	
ORDERNUMBER	1.000000	0.065543	-0.002935	-0.055550	0.0
QUANTITYORDERED	0.065543	1.000000	0.005564	-0.018397	0.0
PRICEEACH	-0.002935	0.005564	1.000000	-0.020965	0.0
ORDERLINENUMBER	-0.055550	-0.018397	-0.020965	1.000000	-0.0
SALES	0.039919	0.551426	0.657841	-0.058400	1.0
QTR_ID	-0.051383	-0.035323	0.008712	0.040716	-0.0
MONTH_ID	-0.039723	-0.039048	0.005152	0.034016	-0.0
YEAR_ID	0.904596	0.069535	-0.005938	-0.057367	0.0
MSRP	-0.010280	0.017881	0.670625	-0.021067	0.0

In [27]: `from sklearn.cluster import KMeans`

In [93]:

```
def WCSS(dataframe):
    wcss_list = []

    for k in range(1,8):


        kmeans_model = KMeans(n_clusters=k)
        kmeans_model.fit(dataframe)
        wcss_value= kmeans_model.inertia_
        wcss_list.append(wcss_value)
        print(f'for k == {k}, wcss is { wcss_value }')
    print("Cluster Centers:",kmeans_model.cluster_centers_)
    print("Feature Names:",kmeans_model.feature_names_in_)
    return wcss_list
```

```
In [98]: list1=WCSS(df2)

for k == 1, wcss is 25407.0000000000022
for k == 2, wcss is 20090.88701217339
for k == 3, wcss is 16909.327212616885
for k == 4, wcss is 14818.002265126062
for k == 5, wcss is 13539.084829579537
for k == 6, wcss is 12546.487375729277
for k == 7, wcss is 11750.52826522585
Cluster Centers: [[-1.65229723e-01 -1.62463528e-01 -1.30802245e+00  5.70264992e-02
 -8.96111961e-01  7.80306529e-01  7.65527784e-01 -4.29486958e-01
 -9.69052753e-01]
 [-8.40236519e-01 -1.78149572e-01  6.28741655e-01 -6.37777720e-04
  1.18321519e-01 -1.02733575e+00 -1.00501458e+00 -3.62202633e-01
  3.66600355e-01]
 [ 1.43907888e+00 -3.35333519e-02 -9.31828018e-02 -9.11582949e-02
 -1.56612779e-01 -1.15395289e+00 -1.14407798e+00  1.69382614e+00
 -1.10099474e-01]
 [-7.53639205e-01 -1.28891156e-02 -1.17312351e+00  5.85510281e-03
 -7.95246500e-01 -1.00106684e+00 -9.85109390e-01 -2.76282170e-01
 -9.75117770e-01]
 [ 6.30810799e-01 -1.96967940e-01  4.70065354e-01  3.94032376e-02
  1.16368961e-02  8.31061956e-01  8.28268814e-01  2.51067374e-01
  1.89649224e-01]
 [ 2.02392884e-01  1.11986731e+00  7.97229433e-01 -2.53219181e-01
  2.03435595e+00 -1.77623653e-02 -2.93408295e-02  2.27266963e-01
  1.27265534e+00]
 [-9.17895108e-01 -6.89458182e-02  6.24701521e-01  1.34017581e-01
  2.59650437e-01  8.46111679e-01  8.32639833e-01 -1.16517009e+00
  4.34142038e-01]]
Feature Names: ['ORDERNUMBER' 'QUANTITYORDERED' 'PRICEEACH' 'ORDERLINENUMBER'
'SALES'
'QTR_ID' 'MONTH_ID' 'YEAR_ID' 'MSRP']
```

```
In [101]: def ElbowCurve(wcss_list):
    k=[1,2,3,4,5,6,7]
    plt.plot(k,wcss_list)
    plt.xlabel('Value of k')
    plt.ylabel('WCSS')
    plt.title('Elbow Curve',loc='right')
```

```
In [102]: ElbowCurve(list1)
```



wcss (within cluster sum of square) >> sum of square of distances of points from the respective centroids
Elbow Graph >> elbow shaped graph that helps us decide the optimal value of k.
Silhouette score >> calculated from silhouette co-efficient. Whichever value of k has highest silhouette score that would be decided for k value.

```
In [39]: from sklearn.metrics import silhouette_score
```



```
In [103]: def SilhoutteScore(dataframe):
    silhouette_score_list = []

    for k in range(2,6):

        kmeans_model_new = KMeans(n_clusters=k)

        y_pred_new = kmeans_model.fit_predict(dataframe)

        silhouette_coefficient = silhouette_score(dataframe,y_pred_new)

        silhouette_score_list.append(silhouette_coefficient)

        print(f'for k == {k},& silhouette score is {silhouette_coefficient}')

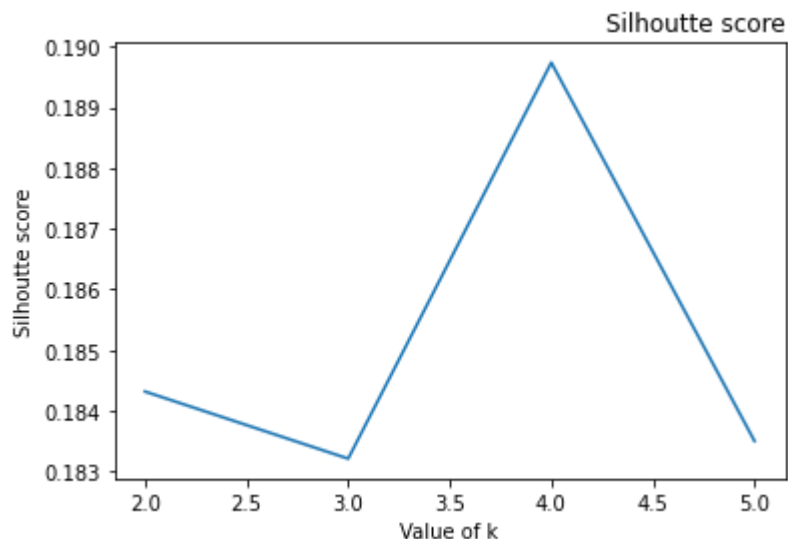
    return silhouette_score_list
```

```
In [104]: slist=SilhoutteScore(df2)
```

```
for k == 2,& silhouette score is 0.18431723406990635
for k == 3,& silhouette score is 0.18321172557165
for k == 4,& silhouette score is 0.18973662495542307
for k == 5,& silhouette score is 0.1835025449952827
```

```
In [105]: def plotSilhoutte(silhouette_score_list):
    k=range(2,6)
    plt.plot(k,silhouette_score_list)
    plt.xlabel('Value of k')
    plt.ylabel('Silhoutte score')
    plt.title('Silhoutte score',loc='right')
```

```
In [106]: plotSilhoutte(slist)
```



```
In [107]: df3=df2[['QUANTITYORDERED', 'SALES']]
```

```
In [89]: df3.shape
```

```
Out[89]: (2823, 2)
```

```
In [90]: df3.sample(5)
```

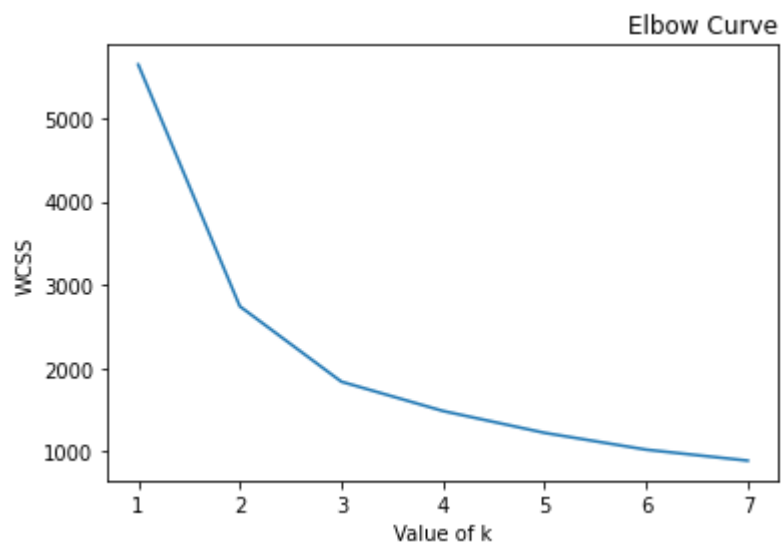
```
Out[90]:
```

	QUANTITYORDERED	SALES
2343	1.530557	0.369860
733	-0.420218	0.509868
2000	-0.728236	-0.923450
2611	0.195816	0.794814
2445	-1.446942	-1.122707

```
In [108]: list2=WCSS(df3)
```

```
for k == 1, wcss is 5646.000000000013
for k == 2, wcss is 2742.1326113529076
for k == 3, wcss is 1836.773942782872
for k == 4, wcss is 1484.9056875574397
for k == 5, wcss is 1222.2810663248342
for k == 6, wcss is 1020.4214771528574
for k == 7, wcss is 888.3489682710209
Cluster Centers: [[-0.9436959 -0.06492512]
 [ 1.06927822 -0.18543648]
 [-1.20680617 -0.98181626]
 [ 0.04988416  0.59361836]
 [-0.06830791 -0.67748743]
 [ 1.69369187  2.92676821]
 [ 0.99929307  1.35167433]]
Feature Names: ['QUANTITYORDERED' 'SALES']
```


```
In [109]: ElbowCurve(list2)
```



```
In [110]: slist1=SilhoutteScore(df3)
```

```
for k == 2,& silhouette score is 0.36926295964297356  
for k == 3,& silhouette score is 0.36925075040605526  
for k == 4,& silhouette score is 0.3693249255482588  
for k == 5,& silhouette score is 0.3692979484278997
```

```
In [111]: plotSilhoutte(slist1)
```



```
In [ ]:
```